

Lecture 8

SQL : Schema Definition, Constraints, and Queries and Views

Relational Database Schema

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

DNUMBER	DLOCATION
---------	-----------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

ESSN	PNO	HOURS
------	-----	-------

DEPENDENT

ESSN	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
------	-----------------------	-----	-------	--------------

Populated Database

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5	
Franklin	T	Wong	333445555	1955-12-08	630 Voss, Houston, TX	M	40000	888665555	5	
Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4	
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4	
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5	
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5	
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4	
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1	

DEPARTMENT	DNAME	DNUMBER	MGRSSN	MGRSTARTDATE	DEPT_LOCATIONS	DNUMBER	DLOCATION
Research	5	333445555	1988-05-22			1	Houston
Administration	4	987654321	1995-01-01			4	Stafford
Headquarters	1	888665555	1981-06-19			5	Bellaire
						5	Sugarland
						5	Houston

WORKS_ON	ESSN	PNO	HOURS
123456789	1	32.5	
123456789	2	7.5	
666884444	3	40.0	
453453453	1	20.0	
453453453	2	20.0	
333445555	2	10.0	
333445555	3	10.0	
333445555	10	10.0	
333445555	20	10.0	
999887777	30	30.0	
999887777	10	10.0	
987987987	10	35.0	
987987987	30	5.0	
987654321	30	20.0	
987654321	20	15.0	
888665555	20	null	

PROJECT	PNAME	PNUMBER	PLOCATION	DNUM
ProductX		1	Bellaire	5
ProductY		2	Sugarland	5
ProductZ		3	Houston	5
Computerization		10	Stafford	4
Reorganization		20	Houston	1
Newbenefits		30	Stafford	4

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATIONSHIP
333445555		Alice	F	1986-04-05	DAUGHTER
333445555		Theodore	M	1983-10-25	SON
333445555		Joy	F	1958-05-03	SPOUSE
987654321		Abner	M	1942-02-28	SPOUSE
123456789		Michael	M	1988-01-04	SON
123456789		Alice	F	1988-12-30	DAUGHTER
123456789		Elizabeth	F	1967-05-05	SPOUSE

The general format of the SELECT statements

- ▶ A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order:

SELECT
FROM
[WHERE
[GROUP BY
[HAVING
[ORDER BY

<attribute list>
<table list>
<condition>]
<grouping attribute(s)>]
<group condition>]
<attribute list>]

Projections (either project all attributes
(*) or project only specific attributes)

Retrieval Queries in SQL

- ▶ **<attribute list>** is a list of attribute names whose values are to be retrieved by the query
- ▶ **<table list>** is a list of the relation names required to process the query
- ▶ **<condition>** is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query
- ▶ We may use the comparison operators :
 $=, <, >, \leq, \geq, \neq$
- ▶ We may use the Boolean operators
AND, OR and NOT

Simple SQL Queries

Example of a simple query on one relation

- ▶ Query 0: Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

Q0: **SELECT BDATE, ADDRESS**
 FROM EMPLOYEE
 WHERE FNAME="John" AND MINIT='B'
 AND LNAME="Smith";

- ▶ The SELECT-clause specifies the projection attributes and the WHERE-clause specifies the selection condition

BDATE	ADDRESS
1965-1-9	171 street

- The result of the query may contain duplicate tuples

UNSPECIFIED WHERE-clause

- ▶ A missing *WHERE-clause* indicates no condition; hence, all tuples of the relations in the *FROM-clause* are selected
 - This is equivalent to the condition WHERE TRUE
- ▶ Query 1: Retrieve the SSN values for all employees.

Q1: **SELECT SSN**
 FROM EMPLOYEE

SSN
12345678
12223334
.....

USE OF *

- ▶ To retrieve all the attribute values of the selected tuples, a * is used, which stands for *all the attributes*

Examples:

Q2: **SELECT ***
 FROM EMPLOYEE
 WHERE DNO=5

USE OF DISTINCT

- ▶ SQL does not treat a relation as a set; duplicate tuples can appear
- ▶ To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used
- ▶ For example, the result of Q3 may have duplicate SALARY values whereas Q3A does not have any duplicate values

Q3 : **SELECT** **SALARY**
 FROM **EMPLOYEE**

Q3A: **SELECT** **DISTINCT** **SALARY**
 FROM **EMPLOYEE**

SALARY
1000
2000
1000
1500
...

SALARY
1000
2000
1500
...

Simple SQL Queries

Query 4: Retrieve the name and address of all employees who work for the 'Research' department.

Q4: **SELECT** **FNAME, LNAME, ADDRESS**
 FROM **EMPLOYEE, DEPARTMENT**
 WHERE **DNAME='Research' AND**
 DNUMBER=DNO

- (DNAME='Research') is a selection condition
- (DNUMBER=DNO) is a join condition

Relational Database Schema

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

DNUMBER	DLOCATION
---------	-----------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

ESSN	PNO	HOURS
------	-----	-------

DEPENDENT

ESSN	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
------	-----------------------	-----	-------	--------------

Aliases, * and DISTINCT, Empty WHERE-clause

- ▶ In SQL, we can use the same name for two (or more) attributes as long as the attributes are in *different relations*
- ▶ A query that refers to two or more attributes with the same name must *qualify* the attribute name with the relation name by *prefixing* the relation name to the attribute name
- ▶ Example:
EMPLOYEE.LNAME, DEPARTMENT.DNAME

Simple SQL Queries

The above query may be re-written using
Relation name.attribute

Q4m: SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME='Research' AND
DEPARTMENT.DNUMBER= EMPLOYEE.DNO

Simple SQL Queries

Query 5: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

Q5:

```
SELECT PNUMBER, DNUM, LNAME, BDATE, ADDRESS
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE DNUM=DNUMBER AND MGRSSN=SSN
AND PLOCATION='Stafford'
```

- In Q5, there are two join conditions
- The join condition DNUM=DNUMBER relates a project to its controlling department
- The join condition MGRSSN=SSN relates the controlling department to the employee who manages that department

ALIASES

- ▶ Some queries need to refer to the same relation twice
 - In this case, *aliases* are given to the relation name
- ▶ Query 6: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

Q6:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM EMPLOYEE E S  
WHERE E.SUPERSSN=S.SSN
```

ALIASES

- In Q6, the alternate relation names E and S are called *aliases* or *tuple variables* for the EMPLOYEE relation
- We can think of E and S as two different *copies* of EMPLOYEE; E represents employees in role of *supervisees* and S represents employees in role of *supervisors*
- ▶ Can use the AS keyword to specify aliases

Q6m:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM   EMPLOYEE AS E, EMPLOYEE AS S  
WHERE  E.SUPERSSN=S.SSN
```

UNSPECIFIED WHERE-clause

- ▶ Example:

**Q7: SELECT SSN, DNAME
FROM EMPLOYEE, DEPARTMENT**

- If more than one relation is specified in the FROM-clause *and* there is no condition, then the *CARTESIAN PRODUCT* of tuples is selected
- It is extremely important not to overlook specifying any selection and join conditions in the WHERE-clause; otherwise, incorrect and very large relations may result

fname	mname	lname	SSN	...									
...													
....													

SET OPERATIONS

- ▶ SQL has directly incorporated some set operations
- ▶ There is a union operation (UNION), and in *some versions* of SQL there are set difference (MINUS) and intersection (INTERSECT) operations
- ▶ The resulting relations of these set operations are sets of tuples; *duplicate tuples are eliminated from the result*
- ▶ The set operations apply only to *union compatible relations*; the two relations must have the same attributes and the attributes must appear in the same order

SET OPERATIONS

Query 8: Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

Q8: (SELECT
 FROM
 WHERE

 UNION

 (SELECT
 FROM
 WHERE

 PNAME
 PROJECT, DEPARTMENT, EMPLOYEE
 DNUM=DNUMBER AND
 MGRSSN=SSN AND LNAME="Smith")

 PNAME
 PROJECT, WORKS_ON, EMPLOYEE
 PNUMBER=PNO AND

 ESSN=SSN AND LNAME="Smith")

Relational Database Schema

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

DNUMBER	DLOCATION
---------	-----------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

ESSN	PNO	HOURS
------	-----	-------

DEPENDENT

ESSN	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
------	-----------------------	-----	-------	--------------

NESTING OF QUERIES

- ▶ A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*
 - ▶ Many of the previous queries can be specified in an alternative form using nesting
- ▶ Query 4: Retrieve the name and address of all employees who work for the 'Research' department.

Q4m: **SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE
WHERE DNO IN**

**(SELECT DNUMBER
FROM DEPARTMENT
WHERE DNAME="Research"**

)

NESTING OF QUERIES

- ▶ The nested query selects the number of the 'Research' department
- ▶ The outer query select an EMPLOYEE tuple if its DNO value is in the result of either nested query
- ▶ The comparison operator **IN** compares a value **v** with a set (or multi-set) of values **V**, and evaluates to TRUE if v is one of the elements in V

NESTING OF QUERIES

- ▶ In general, we can have several levels of nested queries
- ▶ A reference to an *unqualified attribute* refers to the relation declared in the *innermost nested query*
- ▶ In this example, the nested query is *not correlated* with the outer query
- ▶ If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query*, the two queries are said to be *correlated*

CORRELATED NESTED QUERIES

- ▶ Query 9: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
Q9: SELECT          E.FNAME, E.LNAME  
      FROM            EMPLOYEE AS E  
      WHERE           E.SSN IN  
                    (SELECT      ESSN  
                      FROM        DEPENDENT  
                      WHERE       ESSN=E.SSN AND Co-related here 'E.SSN'  
                      E.FNAME=DEPENDENT_NAME)
```

CORRELATED NESTED QUERIES

- ▶ In Q9, the nested query has a different result in the outer query
- ▶ A query written with nested SELECT...
FROM... WHERE... blocks and using the = or
IN comparison operators can *always* be
expressed as a single block query. For
example, Q9 may be written as in Q9A

Nested Queries are best when using aggregate functions. E.g :
Select salary from EMP where SALARY > (select AVG(salary)
from EMP)

Based on Gemini, The performance of one JOIN query is higher than nested query, but it depends on the DB engine, since some engines flatten nested queries to JOIN

Q9A: **SELECT** **E.FNAME, E.LNAME**
FROM **EMPLOYEE E, DEPENDENT D**
WHERE **E.SSN=D.ESSN AND**
 E.FNAME=D.DEPENDENT_NAME

Example

- ▶ Get employee numbers for employees who have the same salary as “John”

```
SELECT SSN
```

```
FROM EMPLOYEE
```

```
WHERE SALARY = ( SELECT SALARY
```

```
    FROM EMPLOYEE
```

```
    WHERE FNAME="John")
```

THE EXISTS FUNCTION

- ▶ EXISTS is used to check whether the result of a correlated nested query is empty (contains no tuples) or not
 - ▶ We can formulate Query 9 in an alternative form that uses EXISTS as Q9B

THE EXISTS FUNCTION

- ▶ Query 9: Retrieve the name of each employee who has a dependent with the same first name as the employee.

```
Q9B: SELECT      FNAME, LNAME  
        FROM       EMPLOYEE  
        WHERE      EXISTS  
                  (SELECT *  
                   FROM    DEPENDENT  
                   WHERE   SSN=ESSN  
                   AND  
                   FNAME=DEPENDENT_NAME)
```

THE EXISTS FUNCTION

- ▶ Query 10: Retrieve the names of employees who have no dependents.

```
Q10:  SELECT  FNAME, LNAME  
        FROM    EMPLOYEE  
        WHERE   NOT EXISTS  
                (SELECT *  
                 FROM  DEPENDENT  
                 WHERE SSN=ESSN)
```

- ▶ In Q10, the correlated nested query retrieves all DEPENDENT tuples related to an EMPLOYEE tuple. If *none exist*, the EMPLOYEE tuple is selected

EXPLICIT SETS

- ▶ It is also possible to use an **explicit (enumerated) set of values** in the WHERE-clause rather than a nested query
- ▶ Query 11: Retrieve the social security numbers of all employees who work on project number 1, 2, or 3.

```
Q11: SELECT DISTINCT ESSN  
      FROM WORKS_ON  
      WHERE PNO IN (1, 2, 3)
```

NULLS IN SQL QUERIES

- ▶ SQL allows queries that check if a value is **NULL** (missing or undefined or not applicable)
- ▶ SQL uses **IS** or **IS NOT** to compare **NULLs** because it considers each **NULL** value distinct from other **NULL** values, so *equality comparison is not appropriate.*
- ▶ Query 12: Retrieve the names of all employees who do not have supervisors.

Q12: **SELECT FNAME, LNAME
 FROM EMPLOYEE
 WHERE SUPERSSN IS NULL**

- Note: If a join condition is specified, tuples with **NULL** values for the join attributes are not included in the result

This is what happened in "Al-Naggar" queries where the query results were missing since joined objects in another table were null.

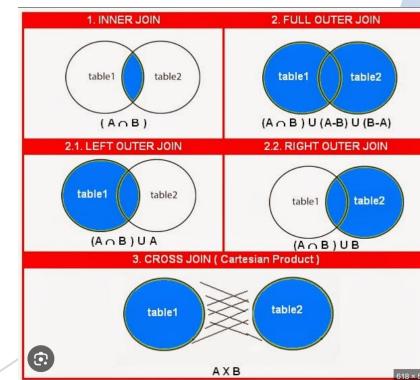
Solution: Use LEFT / RIGHT Joins to include all the table results you need even if null. This is specified in .NET orm with Nullable FK and objects.

Joined Relations Feature in SQL2

- ▶ Can specify a "joined relation" in the FROM-clause
 - Looks like any other relation but is the result of a join
 - Allows the user to specify different types of joins (regular "theta" JOIN, NATURAL JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, CROSS JOIN, etc)

DEFINITIONS

- ▶ **Theta JOIN:** Produces all combinations of tuples from R and S that satisfies the join condition
- ▶ **NATURAL JOIN (inner join) :** As theta join but with equal conditions, and the joined attributes in S are not included
- ▶ **LEFT OUTER JOIN:** R left outer join S means keep all the tuples in R even if they are not matching the conditions.
- ▶ **RIGHT OUTER JOIN:** R right outer join S the same as above but keep all tuples of S



Joined Relations Feature in SQL2

- ▶ Examples:

Q13:

```
SELECT      E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM        EMPLOYEE E S  
WHERE       E.SUPERSSN=S.SSN
```

- ▶ can be written as:

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME  
FROM  (EMPLOYEE E LEFT OUTER JOIN  
        EMPLOYEE S ON E.SUPERSSN=S.SSN)
```

E.FNAME	E.LNAME	S.FNAME	S.LNAME
Hoda	Mohamed	Aly	Ahmed
Aly	Ahmed		

Joined Relations Feature in SQL2

Examples:

Q4: SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME="Research" AND
DNUMBER=DNO

► could be written as:

Q4: SELECT FNAME, LNAME, ADDRESS
FROM (EMPLOYEE JOIN DEPARTMENT
ON DNUMBER=DNO)
WHERE DNAME="Research"

► or as:

Q4: SELECT FNAME, LNAME, ADDRESS
FROM (EMPLOYEE NATURAL JOIN DEPARTMENT
AS DEPT(DNAME, DNO, MSSN, MSDATE))
WHERE DNAME="Research"

Joined Relations Feature in SQL2

- ▶ Another Example: Q5 could be written as follows; this illustrates multiple joins in the joined tables

Q5:

```
SELECT      PNUMBER,DNUM,LNAME,BDATE,ADDRESS  
FROM        (PROJECT JOIN DEPARTMENT ON  
             DNUM=DNUMBER) JOIN EMPLOYEE ON  
             MGRSSN=SSN) )  
WHERE       PLOCATION="Stafford"
```

AGGREGATE FUNCTIONS

- ▶ Include COUNT, SUM, MAX, MIN, and AVG
- ▶ Query 14: Find the maximum salary, the minimum salary, and the average salary among all employees.

Q14: `SELECT MAX(SALARY),
MIN(SALARY), AVG(SALARY)
FROM EMPLOYEE`

- ▶ Some SQL implementations *may not allow more than one function* in the SELECT-clause

AGGREGATE FUNCTIONS

- ▶ Query 15: Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

```
Q15: SELECT MAX(SALARY),  
       MIN(SALARY), AVG(SALARY)  
     FROM EMPLOYEE, DEPARTMENT  
    WHERE DNO=DNUMBER AND  
          DNAME="Research"
```

AGGREGATE FUNCTIONS

- ▶ Queries 16 and 17: Retrieve the total number of employees in the company (Q16), and the number of employees in the 'Research' department (Q17).

Q16: `SELECT COUNT (*)
 FROM EMPLOYEE`

Q17: `SELECT COUNT (*)
 FROM EMPLOYEE, DEPARTMENT
 WHERE DNO=DNUMBER AND
 DNAME="Research"`

GROUPING

- ▶ In many cases, we want to apply the aggregate functions to *subgroups of tuples* in a relation
- ▶ Each subgroup of tuples consists of the set of tuples that have the *same value* for the *grouping attribute(s)*
- ▶ The function is applied to each subgroup independently
- ▶ SQL has a **GROUP BY**-clause for specifying the grouping attributes, which *must also appear in the SELECT-clause*

GROUPING

- ▶ Query 18: For each department, retrieve the department number, the number of employees in the department, and their average salary.

Q18: `SELECT DNO, COUNT (*), AVG (SALARY)
FROM EMPLOYEE
GROUPBY DNO`

Rule, Selecting Columns alongside Aggregate functions must be accompanied with GROUP BY for all columns, otherwise the database engine will not know how to display the selected column (attribute)

- In Q18, the EMPLOYEE tuples are divided into groups-
 - ❖ Each group having the same value for the grouping attribute DNO
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- A join condition can be used in conjunction with grouping

Understanding Grouping: When aggregate functions are used without grouping, it processes all table and have 1 result, but when using GROUP BY, the processing for Aggregate function is across only each group of the specified attribute.
 For e.g , the following table for query having 'GROUP BY DNO' : each set of equal DNO are grouped and processed together with the Aggregate function , displaying aggregate functions for each group independently

SSN	NAME	...	SALARY	DNO
123			1000	1
1234			2000	1
213			100	2
1324			200	2
4356			300	2
12567			500	3
7689			700	4

DNO	COUNT	AVERAGE
1	2	1500
2	3	200
....

GROUPING

- ▶ Query 19: For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
Q19: SELECT PNUMBER, PNAME, COUNT (*)
      FROM PROJECT, WORKS_ON
      WHERE PNUMBER=PNO
      GROUPBY PNUMBER, PNAME
```

- ▶ In this case, the grouping and functions are applied after the joining of the two relations

THE HAVING-CLAUSE

- ▶ Sometimes we want to retrieve the values of these functions for only those *groups that satisfy certain conditions*
- ▶ The HAVING-clause is used for specifying a selection condition on groups (rather than on individual tuples)

'Having' is a 'WHERE' clause applied on Aggregate function (the correct usage) after the grouping to filter the result with the specified condition

THE HAVING-CLAUSE

- ▶ Query 20: For each project *on which more than two employees work*, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT PNUMBER, PNAME, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE PNUMBER=PNO  
GROUPBY PNUMBER, PNAME  
HAVING COUNT (*) > 2
```

PNUMBER	PNAME
1	XX
1	XX
1	XX
2	YY

COUNT=3 > 2

COUNT=1 < 2

PNUMBER	PNAME	COUNT
1	XX	3

SUBSTRING COMPARISON

- ▶ The **LIKE** comparison operator is used to compare partial strings
- ▶ Two reserved characters are used: '%' (or '*' in some implementations) replaces an arbitrary number of characters, and '_' replaces a single arbitrary character

SUBSTRING COMPARISON

- ▶ Query 21: Retrieve all employees whose address is in Houston, Texas. Here, the value of the ADDRESS attribute must contain the substring 'Houston, TX' in it.

Q21: `SELECT FNAME, LNAME
 FROM EMPLOYEE
 WHERE ADDRESS LIKE "%Houston, TX%"`

SUBSTRING COMPARISON

- ▶ Query 22: Retrieve all employees who were born during the 1950s.
 - Here, '5' must be the 8th character of the string (according to our format for date), so the BDATE value is ‘__5____’, with each underscore as a place holder for a single arbitrary character.

Q22: `SELECT FNAME, LNAME
 FROM EMPLOYEE
 WHERE BDATE LIKE "____5_____"`

- ▶ The LIKE operator allows us to get around the fact that each value is considered atomic and indivisible

ARITHMETIC OPERATIONS

- ▶ The standard arithmetic operators '+', '-'. '*', and '/' (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric values in an SQL query result
- ▶ Query 23: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

Q23:

```
SELECT    FNAME, LNAME, 1.1*SALARY  
FROM      EMPLOYEE, WORKS_ON, PROJECT  
WHERE     SSN=ESSN AND PNO=PNUMBER  
          AND PNAME="ProductX"
```

ORDER BY

- ▶ The ORDER BY clause is used to sort the tuples in a query result based on the values of some attribute(s)
- ▶ Query 24: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

Q24: `SELECT DNAME, LNAME, FNAME, PNAME
FROM DEPARTMENT, EMPLOYEE,
 WORKS_ON, PROJECT
WHERE DNUMBER=DNO AND SSN=ESSN
 AND PNO=PNUMBER
ORDER BY DNAME, LNAME`

ORDER BY

- ▶ The default order is in ascending order of values
- ▶ We can specify the keyword **DESC** if we want a descending order; the keyword **ASC** can be used to explicitly specify ascending order, even though it is the default